

Course Project Report

Part II. Neural Networks & Computer Graphics

Pranath Reddy Kumbam
University of Florida, Gainesville, FL 32611, USA

UFID: 8512-0977

In this project, we will be exploring the application of machine learning in computer graphics. Specifically, automatic image colorization. The first part of the project involves the development of a convolutional neural network (CNN)-based regression algorithm for the prediction of mean chrominance values, followed by a colorizer model that colourizes an input image just by taking in the luminance information. The second part of the project is based on a transfer learning approach to develop automatic colorization models for the NCD benchmark dataset.

I. INTRODUCTION

Image colorization is the process of estimating color information for a given grayscale image. In this project, we will be using deep learning algorithms to develop an automatic image colorization pipeline. The application of machine learning for image colorization is an active area of research due to its extensive use in applications such as image and video restoration and also colorization for film and animation.

We will be studying two datasets with LAB color space in this project, specifically, the Georgia Tech Face Dataset and the NCD dataset. In the first part of the project, we will be training and testing regression models for the prediction of mean chrominance values (a^* , b^*) for the Face dataset, followed by an autoencoder-based encoder-decoder model that colorizes images using their luminance values (L). For the second part of the project, we will be using a transfer learning approach to use the models trained on the Faces dataset to develop automatic image colorization models for the NCD dataset. We will be using several methods, such as image augmentation, hyperparameter tuning, and batch normalization, to optimize our implementation.

II. DATA PROCESSING AND IMAGE AUGMENTATION

The Georgia Tech Faces dataset used in the first part of this project consists of images of people and faces. The dataset contains images of 50 people taken in two or three sessions between January 6, 1999, and November 15, 1999, at the Center for Signal and Image Processing at the Georgia Institute of Technology. All people in the database are represented by 15-color JPEG images with cluttered backgrounds, taken at a resolution of 640x480 pixels. The average size of the faces in these images is 150x150 pixels. The pictures show frontal and/or tilted faces with different facial expressions, lighting conditions, and scales. We will be using the OpenCV and NumPy packages in Python to load, process, and augment the images. The images are resized to the dimensions of 128*128 and then divided into training and testing/validation sets using an 80:10:10 split. We then expand the size of the training set by a factor of 10 using image augmentations such as horizontal flips, random crops, and scaling, and we populate the data tensor with the augmented copies using a simple python for-loop. The images are then converted into the LAB color space and saved into separate directories and NumPy arrays which will be used for training and testing the models. We will be using a similar approach to process the NCD dataset.

III. REGRESSION (AND HYPERPARAMETER TUNING)

In this section, I will be discussing the implementation and results of the regression models that predict the mean chrominance values of a given input face sample. I have first trained a simple fully connected model with [8192, 4096, 1024, 256] neurons in each hidden layer for 100 epochs with the learning rate set to 1e-5. The training loss and results are presented in Figure 1 and Table 2 respectively. We then explore convolutional neural network (CNN)-based models. I

| Epochs | LR=1e-3 | LR=1e-4 | LR=1e-5 |
|--------|----------|----------|---------|
| 100 | 0.000196 | 0.000195 | 0.1877 |
| 300 | 0.000183 | 0.000197 | 0.1553 |
| 500 | 0.000198 | 0.000196 | 0.14998 |

(a) CNN1

| Epochs | LR=1e-3 | LR=1e-4 | LR=1e-5 |
|--------|-----------|----------|-----------|
| 100 | 0.0001594 | 0.000128 | 0.000138 |
| 300 | 0.000135 | 0.000128 | 0.0001218 |
| 500 | 0.0001313 | 0.000128 | 0.000108 |

(b) CNN2

| Epochs | LR=1e-3 | LR=1e-4 | LR=1e-5 |
|--------|-----------|----------|----------|
| 100 | 0.000114 | 0.000109 | 0.000142 |
| 300 | 0.000129 | 0.000109 | 0.000136 |
| 500 | 0.0001108 | 0.000138 | 0.000135 |

(c) CNN3

TABLE I: Hyperparameter Tuning, Aggregate MSE Validation Losses of (a^*, b^*)

followed the project description while designing the architecture of the model. The CNN regressor includes 7 convolution blocks with stride and padding adjusted in such a way that the spatial dimension of the image is reduced by a factor of two in each layer. We test three different models with a different set of feature maps. The first model (CNN1) is relatively simple and has three feature maps per hidden layer. The second model (CNN2) has [16, 32, 64, 128, 256, 256] feature maps in hidden layers and the third model (CNN3) has [32, 64, 128, 256, 256, 512] feature maps. I have used grid search for hyperparameter tuning by training and checking the validation loss on various values of learning rate (1e-3, 1e-4, 1e-5) and the number of epochs (100, 300, 500) at a fixed weight decay of 1e-5. The models were trained using the Adam optimizer, and all the results of the validation loss of the tests for hyperparameter tuning are compiled in Table 1. The input values were scaled by dividing by 255, so to analyse the results in the original range, one can simply multiply by 255. As we can see, the CNN2 model performs the best at a learning rate of 1e-5 and 500 epochs. I will be using this architecture and set of hyperparameters for training and testing the final regression model, whose training loss and results are shown in Figure 2 and Table 2 respectively. Additionally, for extra credit, I have also tested the CNN2 model with a Tanh loss, whose training loss and results are shown in Figure 3 and Table 2 respectively. For Tanh's implementation, the values are first scaled to the range [-1,1] before passing to the model and are scaled back while testing.

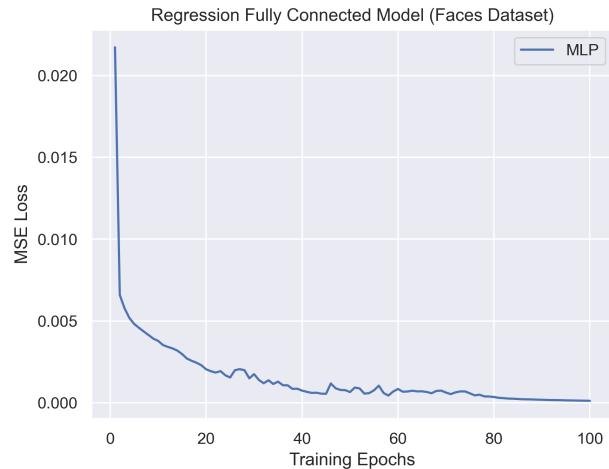


FIG. 1: Training Loss of Fully Connected regression Model

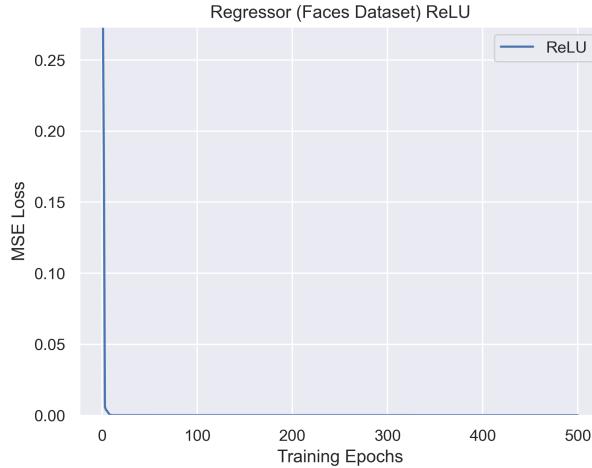


FIG. 2: Training Loss of CNN regression Model with ReLU

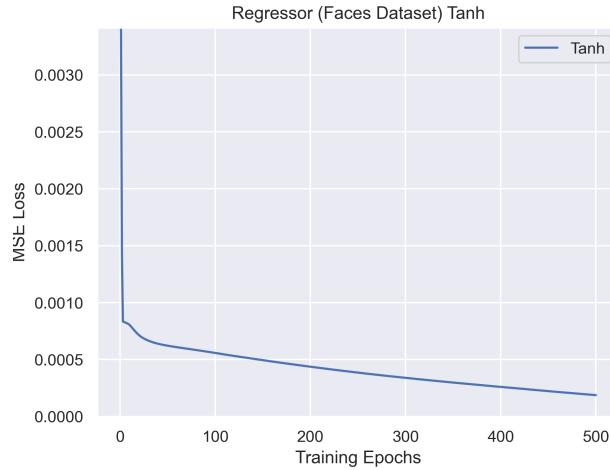


FIG. 3: Training Loss of CNN regression Model with Tanh

TABLE II: Regression MSE Results

| Model | MSE (a*) | MSE (b*) | Learning Rate | Number of Train Epochs |
|------------------------------|-------------|----------|---------------|------------------------|
| Fully Connected Model | 0.000468 | 0.000459 | 1e-5 | 100 |
| CNN with ReLU | 6.2262e-05 | 0.000172 | 1e-5 | 500 |
| CNN with Tanh (Extra Credit) | 4.40178e-05 | 0.000129 | 1e-5 | 500 |

IV. COLORIZATION (AND HYPERPARAMETER TUNING)

We will now be looking at developing encoder-decoder-based models for image colorization by including upsampling or deconvolution layers. The output spatial dimension of the image generated by the network is equal to the input spatial dimensions. As suggested in the project description, I have used five upsampling and downsampling blocks with batch normalization layers added before each spatial convolution layer. For hyper parameter tuning, the model architecture with respect to the number of filters remains the same, but I will be testing three activation functions, i.e., ReLU, Sigmoid, and Tanh for extra credit. Similar to before, I will be using grid

| Epochs | LR=1e-3 | LR=1e-4 | LR=1e-5 |
|--------|----------|----------|----------|
| 100 | 34.5892 | 34.4374 | 34.46079 |
| 300 | 34.12332 | 34.6032 | 34.37195 |
| 500 | 34.57603 | 34.59363 | 34.30788 |

(a) ReLU

| Epochs | LR=1e-3 | LR=1e-4 | LR=1e-5 |
|--------|----------|----------|---------|
| 100 | 34.424 | 34.66826 | 33.6417 |
| 300 | 34.5579 | 34.67858 | 34.6108 |
| 500 | 34.62515 | 34.435 | 34.4634 |

(b) Sigmoid

| Epochs | LR=1e-3 | LR=1e-4 | LR=1e-5 |
|--------|----------|---------|---------|
| 100 | 35.4374 | 35.5301 | 32.1824 |
| 300 | 35.6492 | 35.3373 | 34.9099 |
| 500 | 35.75838 | 35.1595 | 34.8064 |

(c) Tanh (Extra Credit)

TABLE III: Hyperparameter Tuning, PSNE validation results

search for hyperparameter tuning by training and checking the validation loss on various values of learning rate ($1e-3$, $1e-4$, $1e-5$) and the number of epochs (100, 300, 500) at a fixed weight decay of $1e-5$. For Tanh's implementation, the values are first scaled to the range $[-1,1]$ before passing to the model and are scaled back while testing. The models were trained using the Adam optimizer, and all the results of the validation loss of the tests for hyperparameter tuning are compiled in Table 3. As we can see, the best activation is Tanh, followed by Sigmoid, and the best set of learning rates and epochs is 10^{-3} and 500. The metric used for hyperparameter tuning is PSNR, and the final MSE and PSNR results are presented in Table 4. The training loss convergence of the models is shown below, and we also present some of the training and testing colorization samples.

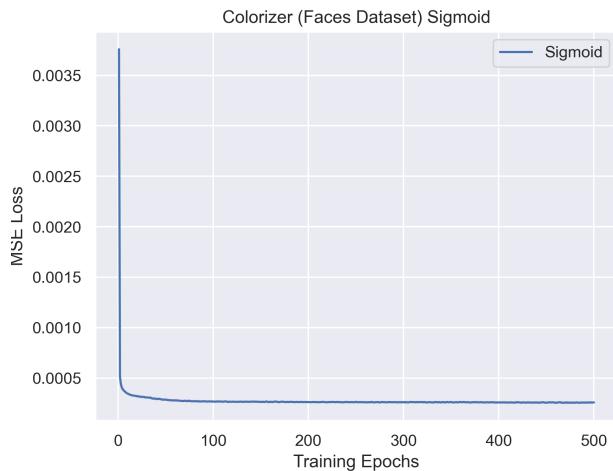


FIG. 4: Training Loss of Colorizer with Sigmoid Activation

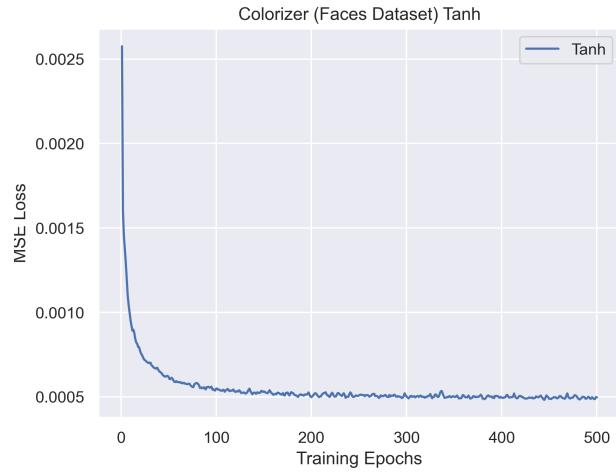


FIG. 5: Training Loss of Colorizer with Tanh Activation

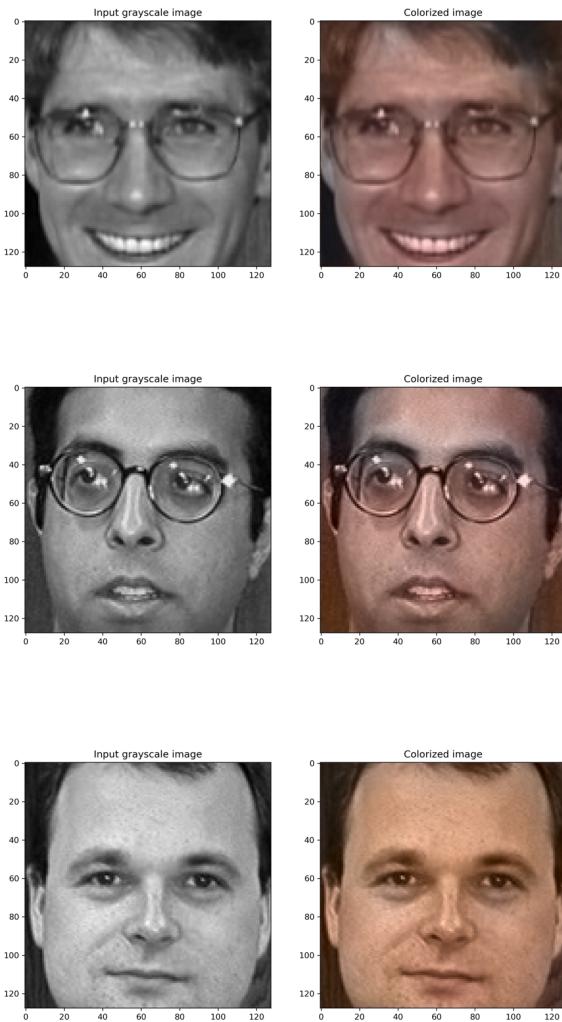


FIG. 6: Example Colorized Train Samples (Sigmoid)

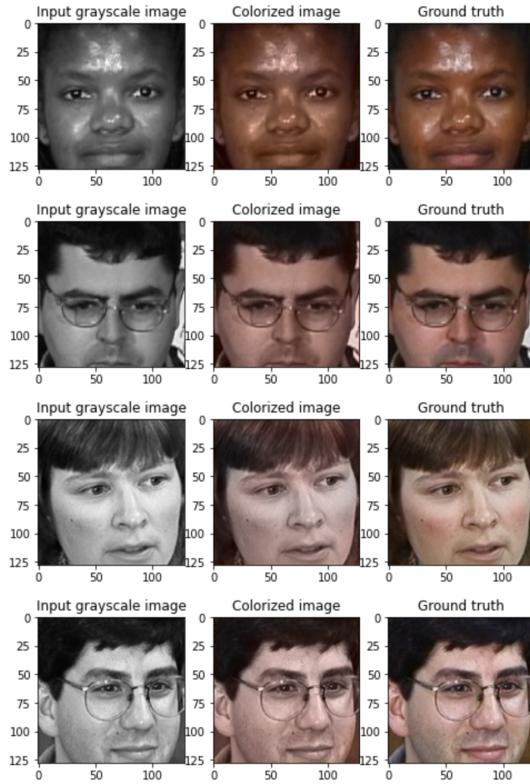


FIG. 7: Example Colorized Test Samples (Sigmoid)

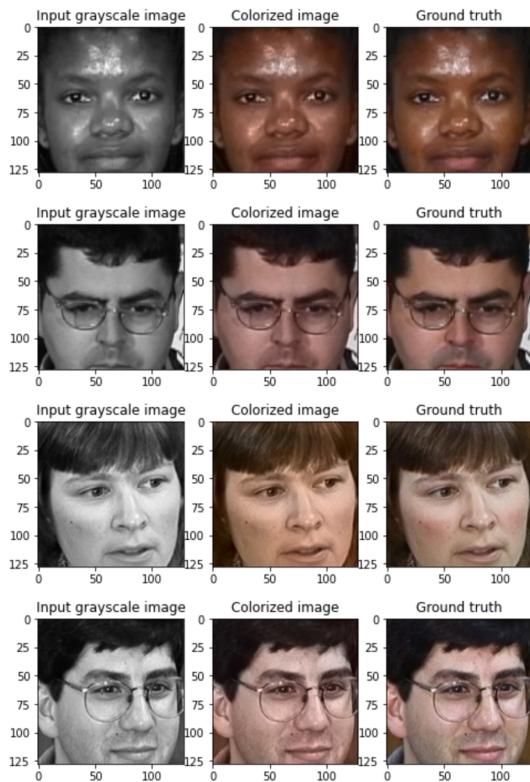


FIG. 8: Example Colorized Test Samples (Tanh)

TABLE IV: Colorization MSE and PSNR Results

| Model | MSE (a*) | MSE (b*) | PSNR (a) | PSNR (b) |
|------------------------------------|----------|----------|----------|----------|
| Colorizer with Sigmoid | 0.000339 | 0.000406 | 34.69402 | 33.90814 |
| Colorizer with Tanh (Extra Credit) | 0.000222 | 0.000380 | 36.5217 | 34.1997 |

V. TRANSFER LEARNING

Next, we apply transfer learning by taking the best-trained models that we got for the Faces dataset and using them for the NCD dataset. Transfer learning is a machine learning technique where a model trained on one task is re-purposed for a second related task. We set the learning rate to 1e-3 and trained the regression and colorizer models that achieved the best scores on the Faces dataset for 1000 epochs on the NCD dataset. I have used ReLU and Tanh (Extra Credit) activation functions for regression and Sigmoid and Tanh (Extra Credit) activation functions for colorization. The training loss plots and the MSE results for regression are shown in Figures 9 and 10 and Table 5 respectively. For colorization, I have used the MSE, PSNR, and SSIM metrics as used in the reference paper. The training loss plots and the results for colorization are shown in Figures 11, 12 and Table 6 respectively. We also show some colorized training and testing samples.

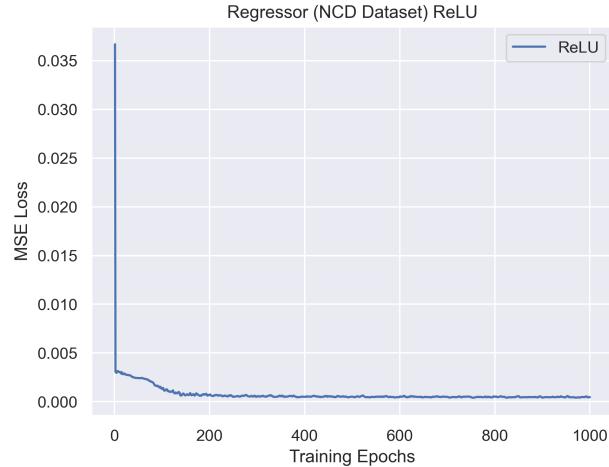


FIG. 9: Training Loss of regression Model with ReLU

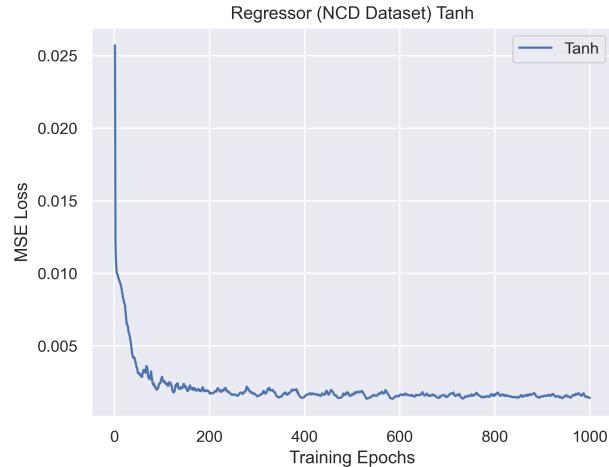


FIG. 10: Training Loss of regression Model with Tanh

TABLE V: Regression MSE Results for NCD

| Model | MSE (a*) | MSE (b*) | Learning Rate | Number of Train Epochs |
|--|----------|----------|---------------|------------------------|
| CNN regressor with ReLU | 0.00134 | 0.00094 | 1e-3 | 1000 |
| CNN regressor with Tanh (Extra Credit) | 0.00164 | 0.00098 | 1e-3 | 1000 |

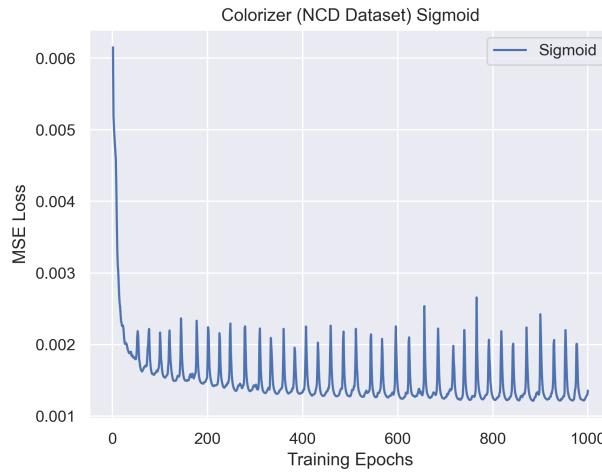


FIG. 11: Training Loss of Colorizer with Sigmoid Activation

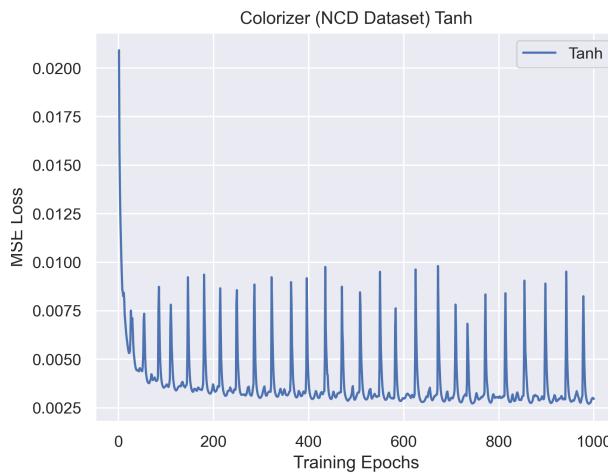


FIG. 12: Training Loss of Colorizer with Tanh Activation

TABLE VI: Colorization MSE, PSNR, and SSIM Results for NCD

| Model | MSE (a*) | MSE (b*) | PSNR (a) | PSNR (b) | SSIM (a) | SSIM (b) |
|------------------------------------|----------|----------|----------|----------|----------|----------|
| Colorizer with Sigmoid | 0.005042 | 0.003354 | 22.97336 | 24.74351 | 0.706159 | 0.80465 |
| Colorizer with Tanh (Extra Credit) | 0.00332 | 0.00299 | 24.7877 | 25.2298 | 0.80334 | 0.83706 |

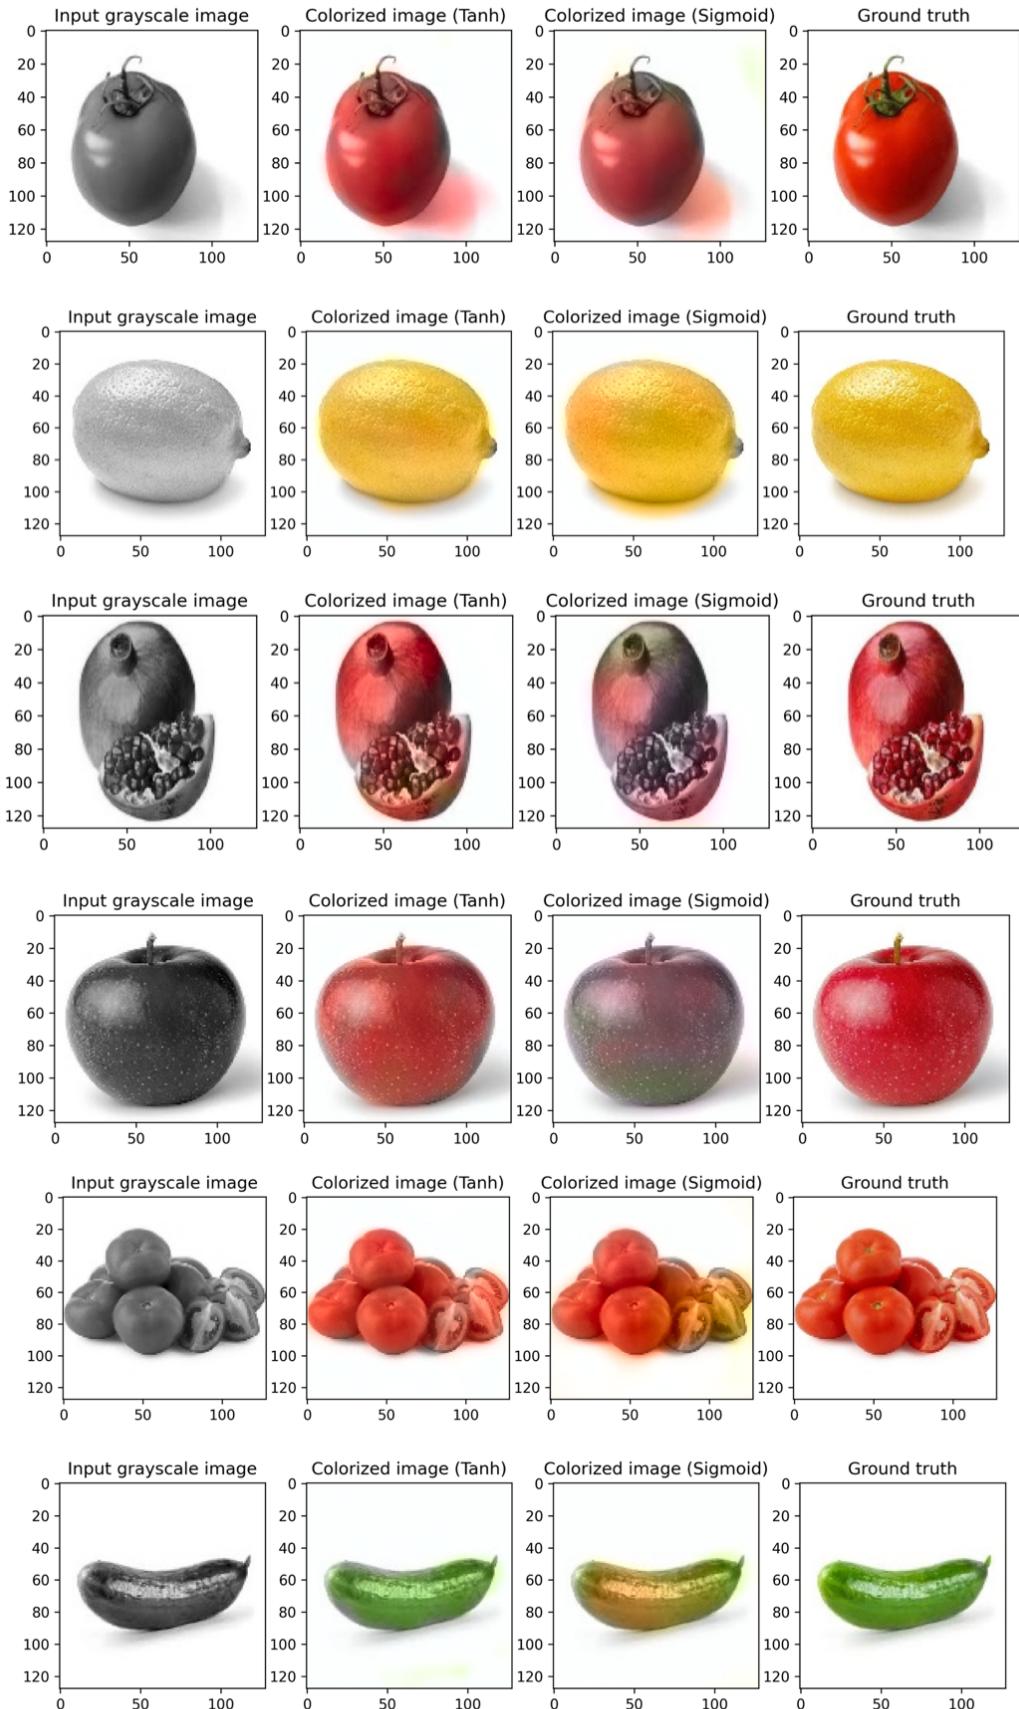


FIG. 13: Selected example Colorized NCD Test Samples (Tanh and Sigmoid)

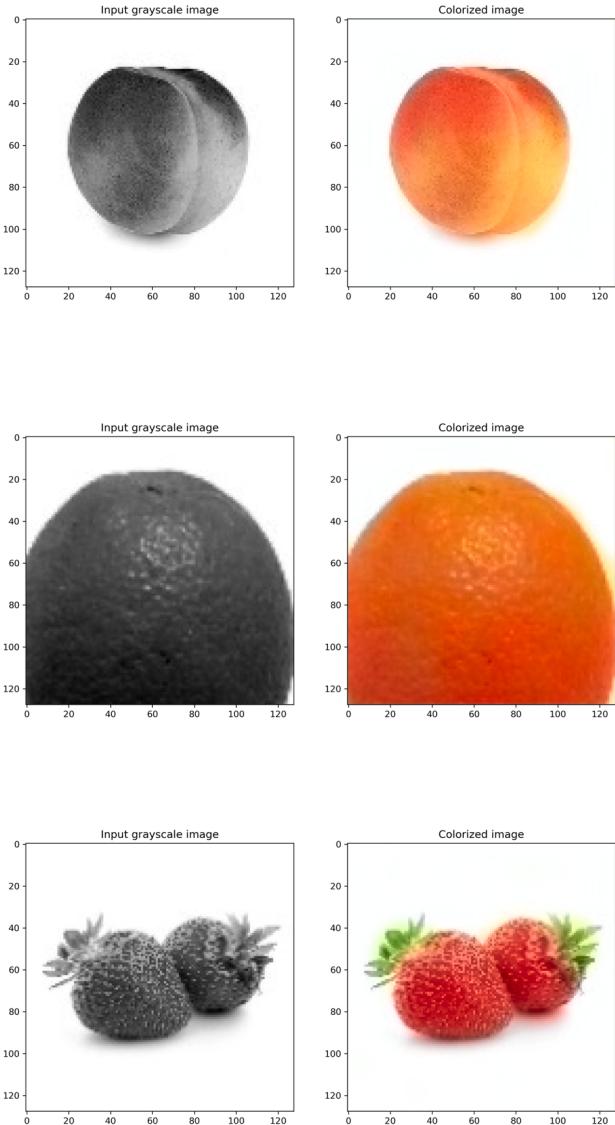


FIG. 14: Example Colorized NCD Train Samples

VI. EXTRA CREDIT

For one of the extra credit components, I have tested the training runtimes on CPU and GPU to compare the runtimes and calculate the speedup we get by moving the model to GPU. I calculated the speedup for a CNN regression model trained for 30 epochs and found the training time to be 8.5 times faster on GPU, which is quite significant. For the other extra credit component, I have already presented all the results with the Tanh activation function in the previous sections and in most cases, Tanh seems to be outperforming equivalent models with ReLU and Sigmoid activations, which is expected since the Tanh function is symmetric about the origin, which leads to easier learning and faster convergence.

VII. SUBMISSION DETAILS

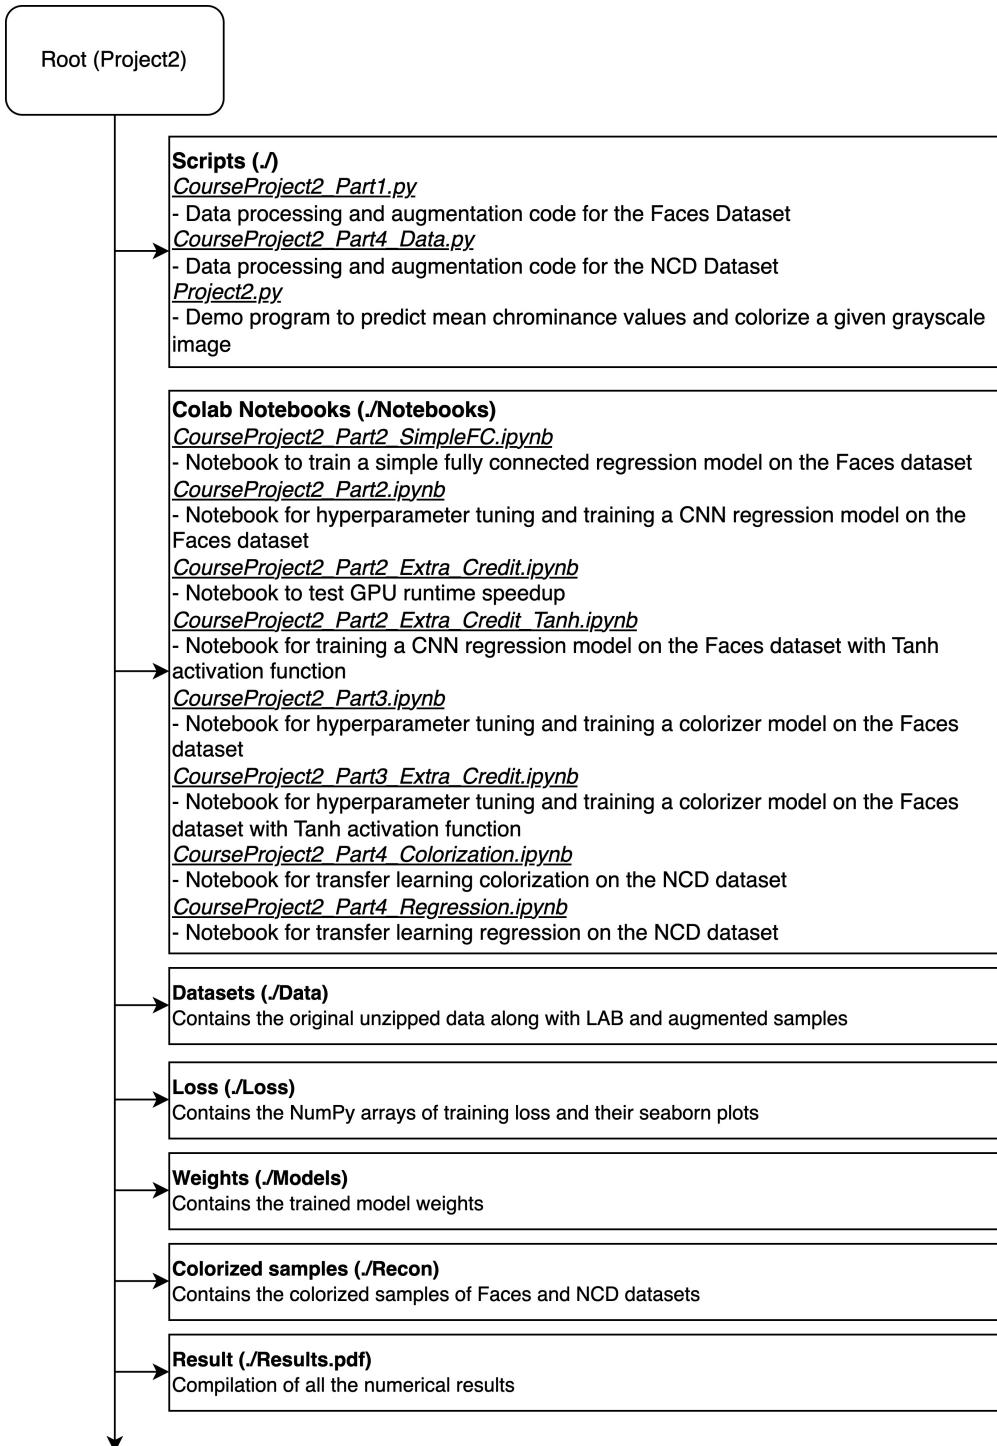


FIG. 15: File structure of the submission

Figure 15 shows the file structure of the submission and the details of each Python script and notebook. The notebooks are divided based on the types of tasks and models. As shown in the file structure, the model weights and data are stored in separate folders. The data folder contains the datasets in LAB format and also the augmented samples. The results are compiled into a pdf file and the colorized samples of Faces and NCD datasets can be found in the Recon folder. The demo script (Project2.py) is a program to predict mean chrominance values and colorize a given grayscale image. The test images can be placed in the './Test' folder and the program prompts the user to set the mode based on the type of test image (Faces or NCD). The folder containing the zipped project folder, file structure, and the project walkthrough video can be found at the following Google drive link:

Google Drive Submission

| Regression | | | | CPU vs GPU (Extra Credit) | Speedup: 8.5X for 30 epochs of training | | | | | | | | |
|---|--|---|--|--|---|---|--|--------------|--------------|----------|-------------------------------------|--|--|
| Simple Fully Connected Model | LR: 1e-5 Epochs: 100 | Weight Decay: 1e-5 | | | | | | | | | | | |
| Result (MSE) | a: 0.000468 | b: 0.000459 | | | | | | | | | | | |
| Colorization | | | | Transfer Learning | | | | | | | | | |
| Hyperparameter Tuning: | Can be done by monitoring validation loss or a simple grid search, we will be performing grid search | | | | | | | | Epochs: 1000 | LR: 1e-3 | Weight Decay: 1e-5, Optimizer: Adam | | |
| Three models: CNN1,2,3 (Different No Of Feature Maps) | | | | Same Model Architecture but different activations | | | | MSE Results | | | | | |
| Grid Search for Epochs: [100, 300, 500] | Learning Rates: [1e-3, 1e-4, 1e-5] | Weight Decay: Fixed at 1e-5 | Adam Optimizer, MSE Loss, ReLU Activation | Grid Search for Epochs: [100, 300, 500] | Learning Rates: [1e-3, 1e-4, 1e-5] | Weight Decay: Fixed at 1e-5 | Adam Optimizer, MSE Loss | Regression | a | | b | | |
| CNN1 | Aggregate MSE Validation Losses of [a,b] | | | | ReLU | PSNR | | | ReLU | 0.00134 | 0.00094 | | |
| Epoch/LR | 1.00E-03 | 1.00E-04 | 1.00E-05 | Epoch/LR | 1.00E-03 | 1.00E-04 | 1.00E-05 | | Tanh | 0.00164 | 0.00098 | | |
| 100 | 0.000196 | 0.000195 | 0.1877 | 100 | 34.5892 | 34.4374 | 34.46079 | Colorization | | | | | |
| 300 | 0.000183 | 0.000197 | 0.1553 | 300 | 34.12332 | 34.6032 | 34.37195 | | | | | | |
| 500 | 0.000198 | 0.000196 | 0.14998 | 500 | 34.57603 | 34.59363 | 34.30788 | Colorization | PSNR Results | | a | | |
| CNN2 | | | | | Sigmoid | | | | Sigmoid | 22.97336 | 24.74351 | | |
| Epoch/LR | 1.00E-03 | 1.00E-04 | 1.00E-05 | Epoch/LR | 1.00E-03 | 1.00E-04 | 1.00E-05 | | Tanh | 24.7877 | 25.2298 | | |
| 100 | 0.0001594 | 0.000128 | 0.000138 | 100 | 34.424 | 34.66826 | 33.6417 | SSIM Results | | | | | |
| 300 | 0.000135 | 0.000128 | 0.0001218 | 300 | 34.5579 | 34.67858 | 34.6108 | a | | b | | | |
| 500 | 0.0001313 | 0.000128 | 0.000108 | 500 | 34.62515 | 34.435 | 34.4634 | Sigmoid | 0.706159 | 0.80465 | | | |
| CNN3 | | | | | Tanh (Extra Credit) | | | | Tanh | 0.80334 | 0.83706 | | |
| Epoch/LR | 1.00E-03 | 1.00E-04 | 1.00E-05 | Epoch/LR | 1.00E-03 | 1.00E-04 | 1.00E-05 | MSE Results | | | | | |
| 100 | 0.000114 | 0.000109 | 0.000142 | 100 | 35.4374 | 35.5301 | 32.1824 | Colorization | a | | b | | |
| 300 | 0.000129 | 0.000109 | 0.000136 | 300 | 35.6492 | 35.3373 | 34.9099 | | Sigmoid | 0.005042 | 0.003354 | | |
| 500 | 0.0001108 | 0.000138 | 0.000135 | 500 | 35.75838 | 35.1595 | 34.8064 | | Tanh | 0.00332 | 0.00299 | | |
| Best Model and hyperparameters: | Model: CNN2 | LR: 1e-5 Epochs: 500 Weight Decay: 1e-5 | This model and set of hyperparameters will be used for transfer learning | Best Model and hyperparameters: | Model: Sigmoid | LR: 1e-3 Epochs: 500 Weight Decay: 1e-5 | This model and set of hyperparameters will be used for transfer learning | | | | | | |
| Result (MSE) | a: 6.2262e-05 | b: 0.000172 | | Result (PSNR) | a: 34.69402 | b: 33.90814 | | | | | | | |
| Testing best model (CNN2) with Tanh (Extra Credit) | a: 4.40178e-05 | b: 0.000129 | Model: CNN2 LR: 1e-5 Epochs: 500 Weight Decay: 1e-5 | Testing best model with Tanh (Extra Credit) (PSNR) | a: 36.5217 | b: 34.1997 | Model: Tanh LR: 1e-3 Epochs: 500 Weight Decay: 1e-5 | | | | | | |
| | | | Result (MSE) | a: 0.000222 | b: 0.000380 | | | | | | | | |